When does data quality help human deferral? Knowledge-based systems for AI governance and managerial decision-making

Yeshwanth Reddy Kunam[a,*]

[a]Midwest University, 851 Parr Road, Wentzville, MO 63385, USA

*Correspondence: 1529 Gage St, Columbus, OH 43240, USA

E-mail: yeshwanth.kunam@gmail.com

Short communication.

Abstract

AI governance must allocate scarce review capacity when data-quality rules prohibit automation. We investigate when quality information improves deferral beyond model probabilities, separating estimation, admission and rule enforcement. A fixed-loss allocator handles mandatory overload exactly; propensity-weighted fitting addresses selective review logs under explicit overlap assumptions. Following exploratory credit experiments, we froze a validation plan before analyzing a new Bank Marketing dataset. In its primary contrast, adding quality only to the automation-loss head reduced cost by 5.23% relative to probability-only estimation under informative missingness (20/20 splits; corrected 95% interval for absolute reduction, 0.0099–0.0224; $p < 0.001$). This supports a conditional benefit, not universal superiority. Across six planned secondary comparisons, inverse-propensity weighting did not significantly improve cost after Holm adjustment. A separate transfer check uses recorded human image classifications without simulating responses to corrupted stimuli. Policy ablations distinguish visible-defect coverage from zero rule violations; service simulations distinguish review priority from capacity under unit service times. The contribution is a reproducible conditions analysis of governance-constrained deferral, with matched component comparisons and an independently evaluated quality hypothesis.

Keywords: Knowledge-based systems; Learning to defer; Data quality; AI governance; Human–AI collaboration; Review capacity; Managerial decision-making

1 Introduction

Organizations divide decisions between models and reviewers. Governance defines controls [1, 2, 3, 4], but oversight alone does not establish effectiveness [5]. Deterioration can increase mandatory review; predictive benefit, rule coverage and capacity differ.

Rejection and learning to defer balance model error against abstention or reviewer error [6, 7, 8, 9]; uncertainty methods offer simpler alternatives [10, 11]. Triage separates model and human error [12]. DeCCaF combines costs, limited feedback and workload constraints [13]; related work addresses sparse feedback [14], quality-aware

cascades [15] and congestion [16]. We test quality under automation restrictions and mandatory holds, claiming no new surrogate or first constrained deferral system.

Completeness and validity incompletely measure quality [17, 18, 19]. Human-in-the-loop and knowledge-based systems connect predictions to action [20, 21, 22]; routing traces document decisions, not model explanations [23, 24, 25]. We ask: *when does quality-aware deferral improve on probability-only estimation and cost-aware heuristics under mandatory review rules?* We provide exact allocation, matched comparisons and planned cross-dataset validation.

2 Method

2.1 Rules and scarce review capacity

A predictor gives p_i and automates with $\hat{y}_i = \mathbf{1}[p_i \geq 1/6]$. False-negative, false-positive, review and hold costs are 5, 1, $c_R = 0.05$ and $c_H = 1.5$. Observable quality $q_i = (u_i, v_i)$ is the missing-field and out-of-range fraction. Training-only bounds are $[Q_{.01,j} - .1t_j, Q_{.99,j} + .1t_j]$, $t_j = \max(Q_{.99,j} - Q_{.01,j}, 1)$. The knowledge base marks $m_i = \mathbf{1}[u_i \geq .5 \vee v_i \geq .25]$. These are explicit experimental policy assumptions, not regulatory thresholds or empirically optimal rules.

Let a_i, r_i, h_i denote automate, review and hold, $M = \{i : m_i = 1\}$, and $B = \lfloor bn \rfloor$. For fixed estimated automation and reviewer losses, allocate by

$$\min \frac{1}{n} \sum_i \left[a_i \hat{\ell}_i^A + r_i(\hat{\ell}_i^H + c_R) + h_i c_H \right], \quad \sum_i h_i = (|M| - B)_+, \tag{1}$$

subject to $a_i + r_i + h_i = 1$, $\sum_i r_i \leq B$, $a_i \leq 1 - m_i$ and $h_i \leq m_i$, with binary actions. The equality assumes mandatory service priority, not an unconstrained economic optimum.

Proposition 1. If $|M| \leq B$, review all mandatory cases and up to $B - |M|$ optional cases with largest positive $\hat{g}_i = \hat{\ell}_i^A - \hat{\ell}_i^H - c_R$. If $|M| > B$, review the B mandatory cases with smallest $\hat{\ell}_i^H$ and hold the rest. *Proof.* Optional review changes cost by $-\hat{g}_i/n$. Under overload, only the sum of reviewed human losses varies; $(|M| - B)c_H$ is constant. Exchange arguments establish both cases. Allocation takes $O(n \log n)$. Exactness is conditional on estimated losses, not a guarantee of minimum realized cost.

Under mandatory overload, changing only the automation-loss head cannot alter routing with a fixed reviewer head. Quality must improve reviewer-loss estimates or policy coverage instead.

Proposition 2. With fixed mandatory flags, identical arrivals, unit review times, no abandonment and work-conserving mandatory service, queue ordering leaves mandatory backlog unchanged. *Proof.* The recurrence $Q_{t+1} =$

$(Q_t + A_t^M - B_t)_+$ depends on counts, not identities. This does not extend automatically to heterogeneous service times, changing flags or reviewer fatigue.

2.2 Loss estimation and selective feedback

With $w(y) = 5y + (1-y)$, targets are $L^A = w(y)\mathbf{1}[\hat{y} \neq y]$ and $L^H = w(y)\mathbf{1}[\text{review} \neq y]$. All gain-fitting rows have outcome labels, but only logged rows have a reviewer decision. Two boosted regressors estimate the losses. Probability inputs are (p, s), $s = \min(5p, 1-p)$; quality inputs add $(u, v, pu, pv, \hat{y}u, \hat{y}v)$. The primary ablation changes only the automation head, retaining the same probability-only reviewer head, admission rule and allocator.

Let S indicate logging with known propensity $\pi(W)$ for observed covariates W. Unweighted fitting minimizes logged-row squared error; inverse-propensity weighting (IPW) uses weights $1/\pi(W)$, normalized to mean one among logged rows [26]. Under $S \perp L^H \mid W$ and $\pi(W) > 0$, $\mathbb{E}[SR_f/\pi(W)] = \mathbb{E}[R_f]$, $R_f = \{L^H - f(Z)\}^2$. This standard identity identifies fitting risk even if Z omits logging covariates; it does not guarantee finite-sample benefit or correct hidden selection, missing outcomes or test shift. ChatGPT and Claude assisted code and analysis development; numerical results were computed with the released Python implementation.

3 Experimental design

3.1 Exploration and prospective validation

The exploratory credit grid uses German, Australian, Default and HMEQ [27, 28, 29, 30, 31], simulated reviewers and nine quality conditions. Adapted surrogate comparisons remain exploratory: no learned router beat the best pooled heuristic, and none of 18 quality ablations survived Holm adjustment.

Before fetching new validation data, we committed the primary contrast, secondary family and complete design (plan `656f527`, implementation `da41f59`). Default supplies development evidence; Bank Marketing supplies a new business-domain validation [32] (45,211 rows). We remove post-call `duration`; categorical mappings use predictor-training rows only, with unknown or unseen categories treated as missing. Historical subscription labels are real, but review outcomes and costs remain simulated. This is cross-domain validation of the routing experiment, not a credit-deployment study.

Feature-identical rows stay together in partitions of 45% training, 15% calibration, 20% gain fitting and 20% testing (seeds 2–21). Histogram gradient boosting (at most 120 iterations, 15 leaves) is logistically calibrated on clean/missing/outlier/mixed copies; gain rows receive the same augmentation. Heads choose three or seven leaves

by grouped 80/20 inner gain validation (100 iterations, rate .06, L2=5), then refit. No augmented copies cross inner folds; test outcomes never select models.

Conditions are clean, informative missingness, stale values and severe defects. Informative missingness multiplies a uniform $[0,.35]$ rate by 1.8 for positives and .6 for negatives. Stale continuous values are multiplied by $[.55,.90]$ on half the rows. Severe shift combines 35–75% missingness and 8–25% outliers. Reviewers have constant error .10 or AI-correlated error, $\min(.45,.05+.25\mathbf{1}[\hat{y}\neq y]+.10[1-|2p_{clean}-1|])$. These mechanisms are illustrative; neither reviewer is defined from quality inputs.

3.2 Logging, matched comparisons and human transfer

Logging is random ($\pi=.4$), uncertainty-prioritized or rule-prioritized. Selective policies use $\pi=.1+.8\sigma(\alpha+z)$ with mean .4, where z is standardized threshold proximity or $4m+2(u+v)$. Shared uniforms and reviewer outcomes give methods identical logs. Comparators are probability-only, quality-in-automation-only, quality-in-both-heads, observed-feature two-head, single-head gain, confidence, Chow and rules-only. The feature comparator includes raw inputs and quality; every head has the same two-option search.

The matched arm shares probability-only reviewer-loss ordering under overload. A fill-budget arm removes optional admission cutoffs. Native, gain-order and no-rule arms isolate allocation and rules at 10% and 20% capacity.

CIFAR-10H supplies recorded human judgments [33]; ten-class choices are mapped to animal/vehicle. The predictor uses 8×8 averaged RGB features on 10,000 original images. Gain fitting samples one recorded choice per image; held-out empirical human error supplies evaluation loss. No corruption is applied: people saw clean images. Logging, head search and 20% capacity are unchanged. This is neither finance-expert validation nor a test of the primary quality hypothesis.

3.3 Endpoints and inference

The single primary test compares quality-in-automation-only with probability-only on Bank under informative missingness, random logging and 20% capacity, averaging reviewers per split. Success requires positive reduction, two-sided $p<.05$ and at least 1% relative improvement. Corrected resampled tests use ratio $.2/.8$ [34]; splits are dependent. Six secondary tests compare IPW versus unweighted full-quality fitting for two datasets and three logging policies, averaging conditions/reviewers at 20%, with one Holm adjustment [35]. Other analyses are exploratory; ambiguous earlier seven-/nine-test claims are retired.

4 Results and discussion

A conditional quality benefit was confirmed. The planned contrast reduced cost by 5.23% (absolute reduction 0.0161; corrected 95% interval [0.0099, 0.0224]; $p < .001$; 20/20 splits), meeting the 1% practical criterion (Table 2). Reviewer estimation and mandatory ordering are fixed; the fill-budget arm retains the same reduction. Figure 1 reports all condition/logging cells. Under mandatory overload, the automation-head effect is zero, as predicted. Confirmation concerns this synthetic mechanism on new data, not universal superiority.

Matched comparisons bound the claim. Table 1 uses a common policy. Quality in the automation head had the lowest mean among the matched variants on Default (0.5966) and Bank (0.4663); Chow cost risk gave 0.6018 and 0.4685, respectively. Lower means do not establish cross-setting significance.

Selective feedback did not make IPW uniformly useful. None of six secondary contrasts survives Holm adjustment (Table 2). Rule-prioritized logging selects approximately 89% of mandatory gain rows versus 40% under random logging at the same expected budget. Variance, shift and unchanged review selection can offset weighting benefits. Human-transfer results are descriptive: on clean images under random logging, quality-in-automation-only cost was 0.1615, probability-only 0.1615, and Chow 0.1554; these means do not establish a human-domain quality effect.

Table 1: Expected cost under random logging and 20% capacity. Default/Bank average four conditions and two reviewers; CIFAR-10H uses clean images and empirical human error. Common mandatory ordering. Bold: lowest descriptive mean, not significance.

Method	Default	Bank	CIFAR-10H
Rules only	0.6703	0.5240	0.2778
Confidence	0.6039	0.4725	0.1670
Chow cost risk	0.6018	0.4685	**0.1554**
Single-head gain	0.6000	0.4682	0.1606
Probability-only two-head	0.6043	0.4706	0.1615
Quality in automation head	**0.5966**	**0.4663**	0.1615
Quality in both heads	0.5981	0.4674	0.1615
Observed-feature two-head	0.6030	0.4669	0.1579

Knowledge rules have costs and blind spots. Across Default and Bank, rules flagged 67.1% of records with

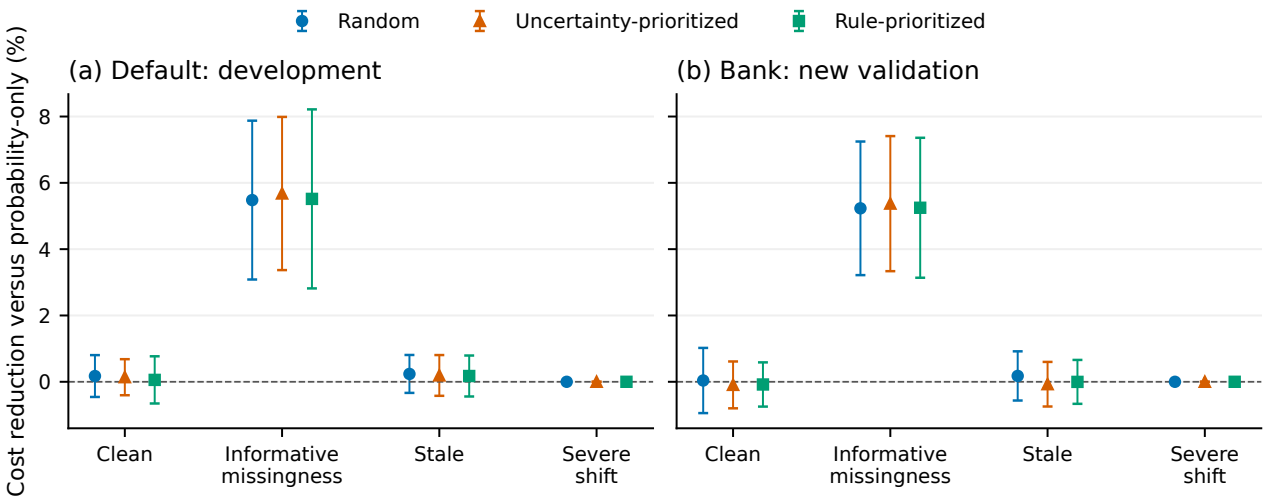


Figure 1: Quality in the automation head versus probability-only at 20% capacity, with common mandatory ordering. Positive reduction favors quality. Bars: approximate corrected 95% intervals over 20 splits, averaging reviewers. Only Bank/informative/random is confirmatory. Programmatically generated from saved outputs.

Table 2: Absolute cost reduction (positive favors the change), corrected 95% interval and p-value. The primary test stands alone; secondary p-values are Holm-adjusted across six contrasts.

Contrast	Reduction	95% interval	p
Primary: Bank, quality in automation head	+0.0161	[+0.0099, +0.0224]	< 0.001
IPW: Default, random	+0.0000	[+0.0000, +0.0000]	1.000
IPW: Default, uncertainty logs	+0.0007	[-0.0002, +0.0017]	0.641
IPW: Default, rule-prioritized logs	-0.0003	[-0.0018, +0.0012]	1.000
IPW: Bank, random	+0.0000	[+0.0000, +0.0000]	1.000
IPW: Bank, uncertainty logs	+0.0000	[-0.0008, +0.0008]	1.000
IPW: Bank, rule-prioritized logs	+0.0001	[-0.0009, +0.0011]	1.000

injected severe defects but 0.0% with injected stale values. Under severe shift, the combined router held 47.1% of all records; removing rules automated 51.9% of all records despite mandatory flags. Zero violations do not establish defect coverage or legal sufficiency; stale but plausible values can escape the rules.

Capacity results require their assumptions. The 40-day simulation introduces severe defects on days 11–20 and carries holds forward. Service counts all eligible mandatory arrivals, with unresolved cases unmet; quantiles account for censoring [36]. Cost ordering leaves 21.7% open at 20% capacity, none at 40%. At 40%, FIFO lowers

the 95th-percentile delay from 13.9 to 7.1 days despite identical backlog counts (Fig. 2). This depends on unit service and fixed flags.

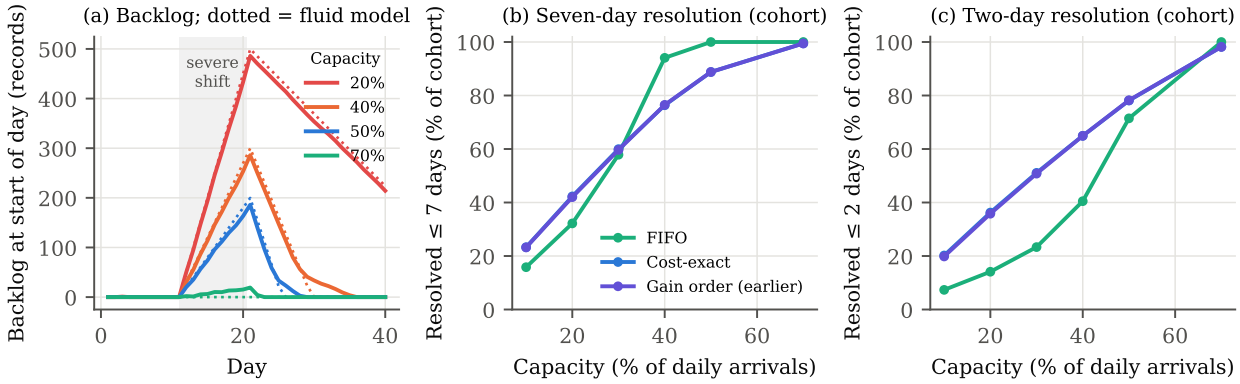


Figure 2: Credit simulation: (a) backlog, with dotted fluid approximation; (b,c) seven-/two-day resolution among eligible mandatory arrivals, unresolved cases unmet. Four datasets, two reviewers, ten seeds; unit service times. Programmatically generated from saved outputs.

Managers should benchmark against local cost-aware heuristics, document rules, maintain exploration under selective logging and monitor unresolved cohorts. Organizational complementarity [37] and AI capability [38] motivate these choices; causal firm benefits are untested.

5 Limitations and conclusion

Bank is old and synthetically corrupted. Human images do not validate finance workflows, damaged stimuli or fatigue. Costs and thresholds need domain validation; gradual drift [39], heterogeneous service, fairness and missing outcomes remain untested. IPW does not correct hidden selection. Surrogate adaptations remain exploratory; no newly tuned DeCCaF comparison or new learning guarantee is claimed.

A frozen validation supports a conditional quality benefit, null weighting results and no universal superiority. Separating estimation, rules and capacity makes these boundaries observable.

Data and code availability Code, outputs, the frozen plan and source hashes: public snapshot `db2ab09`. Original credit evidence: `03166e9`. Source data: cited providers. Human-derived outputs follow CC BY-NC-SA 4.0; original images and participant-level records are not redistributed here.

Funding This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

CRediT authorship contribution statement Yeshwanth Reddy Kunam: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Writing – original draft, Writing – review and editing, Visualization.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process ChatGPT (OpenAI) assisted study design, code, drafting, experiments and verification; Claude (Anthropic) assisted earlier experimental extensions, analysis and drafting. Figures were generated from saved outputs. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the content of the published article.

References

[1] R. Abraham, J. Schneider, J. vom Brocke, Data governance: A conceptual framework, structured review, and research agenda, Int. J. Inf. Manag. 49 (2019) 424–438. doi:10.1016/j.ijinfomgt.2019.07.008.

[2] M. Janssen, P. Brous, E. Estevez, L.S. Barbosa, T. Janowski, Data governance: Organizing data for trustworthy artificial intelligence, Gov. Inf. Q. 37 (2020) 101493. doi:10.1016/j.giq.2020.101493.

[3] E. Papagiannidis, P. Mikalef, K. Conboy, Responsible artificial intelligence governance: A review and research framework, J. Strateg. Inf. Syst. 34 (2025) 101885. doi:10.1016/j.jsis.2024.101885.

[4] B.W. Wirtz, J.C. Weyerer, I. Kehl, Governance of artificial intelligence: A risk and guideline-based integrative framework, Gov. Inf. Q. 39 (2022) 101685. doi:10.1016/j.giq.2022.101685.

[5] B. Green, The flaws of policies requiring human oversight of government algorithms, Comput. Law Secur. Rev. 45 (2022) 105681. doi:10.1016/j.clsr.2022.105681.

[6] C. Chow, On optimum recognition error and reject tradeoff, IEEE Trans. Inf. Theory 16 (1970) 41–46. doi:10.1109/TIT.1970.1054406.

[7] X.Y. Zhang, G.S. Xie, X. Li, T. Mei, C.L. Liu, A survey on learning to reject, Proc. IEEE 111 (2023) 185–215. doi:10.1109/JPROC.2023.3238024.

[8] H. Mozannar, D. Sontag, Consistent estimators for learning to defer to an expert, in: Proc. 37th Int. Conf. Mach. Learn., PMLR 119 (2020) 7076–7087.

[9] R. Verma, E. Nalisnick, Calibrated learning to defer with one-vs-all classifiers, in: Proc. 39th Int. Conf. Mach. Learn., PMLR 162 (2022) 22184–22202.

[10] J. Mena, O. Pujol, J. Vitrià, Uncertainty-based rejection wrappers for black-box classifiers, IEEE Access 8 (2020) 101721–101746. doi:10.1109/ACCESS.2020.2996495.

[11] M. Abdar, F. Pourpanah, S. Hussain, et al., A review of uncertainty quantification in deep learning: Techniques, applications and challenges, Inf. Fusion 76 (2021) 243–297. doi:10.1016/j.inffus.2021.05.008.

[12] M. Raghu, K. Blumer, G. Corrado, J. Kleinberg, Z. Obermeyer, S. Mullainathan, The algorithmic automation problem: Prediction, triage, and human effort, arXiv preprint (2019). https://doi.org/10.48550/arXiv.1903.12220.

[13] J.V. Alves, D. Leitão, S. Jesus, M.O.P. Sampaio, J. Liébana, P. Saleiro, M.A.T. Figueiredo, P. Bizarro, Cost-sensitive learning to defer to multiple experts with workload constraints, Trans. Mach. Learn. Res. (2024). https://openreview.net/forum?id=TAvGZm2Rqb.

[14] P. Hemmer, L. Thede, M. Vössing, J. Jakubik, N. Kühl, Learning to defer with limited expert predictions, Proc. AAAI Conf. Artif. Intell. 37 (2023) 6002–6011. doi:10.1609/aaai.v37i5.25742.

[15] A. Farinhas, N.M. Guerreiro, S. Agrawal, R. Rei, A. Martins, Translate smart, not hard: Cascaded translation systems with quality-aware deferral, in: Proc. 2025 Conf. Empirical Methods Nat. Lang. Process. (2025) 26742–26756. doi:10.18653/v1/2025.emnlp-main.1358.

[16] V. Saketos, W. Reveillard, A. Proutiere, S. Dalleiger, M. Xiao, To defer or not to defer? Online learning with LLMs and congested experts, arXiv preprint (2026). https://doi.org/10.48550/arXiv.2610.15128.

[17] C. Chai, J. Wang, Y. Luo, Z. Niu, G. Li, Data management for machine learning: A survey, IEEE Trans. Knowl. Data Eng. 35 (2023) 4646–4667. doi:10.1109/TKDE.2022.3148237.

[18] L. Zhang, D. Jeong, S. Lee, Data quality management in the Internet of Things, Sensors 21 (2021) 5834. doi:10.3390/s21175834.

[19] S. Mohammed, L. Budach, M. Feuerpfeil, N. Ihde, A. Nathansen, N. Noack, H. Patzlaff, F. Naumann, H. Harmouch, The effects of data quality on machine learning performance on tabular data, Inf. Syst. 132 (2025) 102549. doi:10.1016/j.is.2025.102549.

[20] X. Wu, L. Xiao, Y. Sun, J. Zhang, T. Ma, L. He, A survey of human-in-the-loop for machine learning, Future Gener. Comput. Syst. 135 (2022) 364–381. doi:10.1016/j.future.2022.05.014.

[21] S. Tariq, M. Baruwal Chhetri, S. Nepal, C. Paris, A^2C: A modular multi-stage collaborative decision framework for human–AI teams, Expert Syst. Appl. 282 (2025) 127318. doi:10.1016/j.eswa.2025.127318.

[22] M. Abbaspour Onari, M. Jahangoshai Rezaee, M. Saberi, M.S. Nobile, An explainable data-driven decision support framework for strategic customer development, Knowl.-Based Syst. 295 (2024) 111761. doi:10.1016/j.knosys.2024.111761.

[23] A. Barredo Arrieta, N. Díaz-Rodríguez, J. Del Ser, et al., Explainable artificial intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI, Inf. Fusion 58 (2020) 82–115. doi:10.1016/j.inffus.2019.12.012.

[24] A. Adadi, M. Berrada, Peeking inside the black-box: A survey on explainable artificial intelligence (XAI), IEEE Access 6 (2018) 52138–52160. doi:10.1109/ACCESS.2018.2870052.

[25] A. Artelt, R. Visser, B. Hammer, "I do not know! but why?" — Local model-agnostic example-based explanations of reject, Neurocomputing 558 (2023) 126722. doi:10.1016/j.neucom.2023.126722.

[26] D.G. Horvitz, D.J. Thompson, A generalization of sampling without replacement from a finite universe, J. Am. Stat. Assoc. 47 (1952) 663–685. doi:10.1080/01621459.1952.10483446.

[27] S. Lessmann, B. Baesens, H.V. Seow, L.C. Thomas, Benchmarking state-of-the-art classification algorithms for credit scoring: An update of research, Eur. J. Oper. Res. 247 (2015) 124–136. doi:10.1016/j.ejor.2015.05.030.

[28] [dataset] H. Hofmann, Statlog (German Credit Data), UCI Machine Learning Repository (1994). https://doi.org/10.24432/C5NC77.

[29] [dataset] R. Quinlan, Statlog (Australian Credit Approval), UCI Machine Learning Repository (1987). https://doi.org/10.24432/C59012.

[30] [dataset] I. Yeh, Default of Credit Card Clients, UCI Machine Learning Repository (2009). https://doi.org/10.24432/C55S3H.

[31] [dataset] B. Baesens, D. Rösch, H. Scheule, Credit Risk Analytics: Measurement Techniques, Applications, and Examples in SAS, Wiley, Hoboken, NJ, 2016 (HMEQ data set; companion data: https://www.creditriskanalytics.net/datasets-private2.html, accessed 23 October 2026).

[32] [dataset] S. Moro, P. Rita, P. Cortez, Bank Marketing, UCI Machine Learning Repository (2014). https://doi.org/10.24432/C5K306.

[33] J.C. Peterson, R.M. Battleday, T.L. Griffiths, O. Russakovsky, Human uncertainty makes classification more robust, in: Proc. IEEE/CVF Int. Conf. Comput. Vis. (2019) 9617–9626. https://openaccess.thecvf.com/content_ICCV_2019/html/Peterson_Human_Uncertainty_Makes_Classification_More_Robust_ICCV_2019_paper.html.

[34] C. Nadeau, Y. Bengio, Inference for the generalization error, Mach. Learn. 52 (2003) 239–281. doi:10.1023/A:1024068626366.

[35] S. Holm, A simple sequentially rejective multiple test procedure, Scand. J. Stat. 6 (1979) 65–70.

[36] E.L. Kaplan, P. Meier, Nonparametric estimation from incomplete observations, J. Am. Stat. Assoc. 53 (1958) 457–481. doi:10.1080/01621459.1958.10501452.

[37] M.H. Jarrahi, Artificial intelligence and the future of work: Human-AI symbiosis in organizational decision making, Bus. Horiz. 61 (2018) 577–586. doi:10.1016/j.bushor.2018.03.007.

[38] P. Mikalef, M. Gupta, Artificial intelligence capability: Conceptualization, measurement calibration, and empirical study on its impact on organizational creativity and firm performance, Inf. Manag. 58 (2021) 103434. doi:10.1016/j.im.2021.103434.

[39] J. Lu, A. Liu, F. Dong, F. Gu, J. Gama, G. Zhang, Learning under concept drift: A review, IEEE Trans. Knowl. Data Eng. 31 (2019) 2346–2363. doi:10.1109/TKDE.2018.2876857.